

NARESH LUHAR

GenAI Engineer

malviyanaresh6219@gmail.com • +91 7357027318 • Mumbai, India • linkedin.com/in/naresh-luhar-8121a6205

PROFESSIONAL SUMMARY

GenAI Engineer with 3+ years of experience architecting production-grade multi-agent systems, enterprise RAG platforms, and real-time voice AI that have cut manual processing time by up to 80% and support automation time by up to 70% across fintech, real estate, manufacturing, and healthcare. Deep expertise in LangGraph, LangChain, GPT-4, and LLaMA, with a track record of taking systems from architecture through evaluation (RAGAS, Promptfoo) to production deployment. Skilled at translating ambiguous business problems into agentic workflows that measurably reduce manual effort and operational cost.

TECHNICAL SKILLS

LLMs & Frameworks: GPT-4, LLaMA, LangChain, LangGraph, CrewAI, Agno, HuggingFace, Ollama, Pipecat, LiveKit

Evaluation & Testing: RAGAS, Promptfoo

AI Specializations: RAG Pipelines, Agentic AI, Multi-Agent Systems, Model Context Protocol (MCP), Prompt Engineering, Tool Calling, Function Calling, Multimodal AI, AI Workflow Orchestration, LoRA Fine-Tuning

Computer Vision & NLP: TensorFlow, PyTorch, BERT, Transformer, ResNet50, STT/TTS, Sentiment Analysis

Data & Databases: PostgreSQL, FAISS, ChromaDB, Vector Databases, DBeaver, Data Analysis, Web Scraping (Selenium)

APIs & Cloud: FastAPI, REST API Integration, Microsoft Azure AI, Gupshup API, MCP Systems

Languages & Tools: Python, SQL, Git, Docker

PROFESSIONAL EXPERIENCE

GenAI Engineer | Logicloop

Sep 2025 Jun 2026

Email AI Agent Mswipe (Fintech)

- Built an AI-powered email agent automating customer complaint processing for payment, KYC, and device issues — reducing manual support ticket creation time by an estimated 70%
- Developed a multimodal extraction pipeline processing emails, attachments, and threads using OCR and LLMs to extract customer, device, and issue details
- Integrated with helpdesk APIs for automated ticket creation and designed missing-information workflows with automated email responses

FactoryGuard AI Enterprise Multi-Agent Predictive Maintenance Platform (Manufacturing)

- Architected a LangGraph-based multi-agent system for predictive maintenance, orchestrating a Health Monitoring Agent for real-time CNC sensor analysis, an RCA Agent for RAG-driven failure diagnosis, and a Maintenance Planning Agent for automated repair recommendations
- Designed a dual-RAG architecture over machine manuals, troubleshooting guides, historical failure records, and maintenance SOPs to deliver explainable root-cause analysis and automate maintenance workflows, reducing unplanned machine downtime

Real-Time Voice AI Agent K Raheja Residency (Real Estate)

- Built a low-latency Voice AI Agent using Pipecat and LiveKit, implementing a full streaming STT → RAG → OpenAI → TTS pipeline for natural real-time conversations
- Designed real-estate knowledge retrieval workflows to answer live queries on flats, pricing, amenities, and locations; stored call transcripts for business analytics

Enterprise Agentic RAG Platform

- Architected a modular enterprise RAG platform for document intelligence and knowledge retrieval, integrating Docling and Unstructured for advanced PDF/enterprise document parsing
- Built end-to-end chunking, embedding, and vector search pipelines using HuggingFace, Ollama, and PostgreSQL; designed reusable agentic RAG ingestion and retrieval workflows
- Evaluated RAG pipeline performance using RAGAS and Promptfoo to measure answer quality, faithfulness, and retrieval effectiveness, enabling iterative prompt and retrieval optimization

AI Backend Interface SolidX MCP System

- Built an MCP-powered AI interface enabling natural-language interaction with a low-code platform, exposing capabilities as LLM-callable tools

- Implemented tool-calling workflows that translated user requests into backend API operations, automating platform configuration and application creation

OCR Intelligence System Reliance

- Built an AI digitization system for handwritten customer forms from fuel stations (name, phone, vehicle details, DOB) using Qwen2.5-VL with advanced prompt engineering
- Evaluated and benchmarked multiple OCR and vision-language models, selecting an optimized open-source solution achieving high accuracy on noisy, inconsistent handwriting

Custom LLM Fine-Tuning with LoRA Enterprise Document Q&A

- Fine-tuned an open-source LLM using Hugging Face Transformers and PEFT (LoRA) on a custom dataset of 3,000+ instruction-response (Q&A) pairs created from enterprise documents, improving domain-specific response accuracy and reducing hallucinations

AI Researcher | Konverge Technologies Mar 2024 May 2025

- Built a GenAI-powered financial RAG system using LangChain, GPT-4, and FAISS to extract real-time insights from financial news via natural language queries
- Engineered a medical document intelligence pipeline using LLaMA and OCR to extract structured data from medical bills, reducing manual data entry by ~80%
- Developed and fine-tuned a ResNet50-based breast X-ray classification model achieving 91%+ accuracy for early-stage anomaly detection
- Built an OCR-driven analytics system processing garage bills to identify frequently replaced components, supporting inventory planning and demand forecasting
- Developed a domain-specific educational chatbot and no-code Q&A framework using TF-IDF for rapid deployment across multiple domains

AI Developer | Cloudstrats Technologies Jun 2022 Dec 2023

- Built NLP systems for multilingual text translation across 10+ Indian languages using LSTM-based models
- Developed GPT-2-based solutions for text generation and abstractive summarization; built BERT-based sentiment analysis engine for large-scale social media data
- Designed a Hindi speech-to-text system with noise reduction for real-world audio environments
- Engineered intelligent document processing solutions to extract structured data from identity documents (Aadhaar, DL, passports, bills)
- Built automated Selenium-based data scraping pipelines to curate datasets for computer vision applications

SELECTED PROJECTS

ATOM Personalized Health Chatbot

- Conversational AI assistant for health-related queries using a domain-specific RAG pipeline with LLM integration and memory management for context-based responses

Music Recommendation System

- Built a music suggestion engine using collaborative filtering, recommending songs based on user interaction history and mood-based clustering

EDUCATION & CERTIFICATIONS

Certification in Data Science and AI | ITvedant Institute, Mumbai Oct 2021 May 2022

Bachelor of Computer Applications | Bhupal Nobles University, Udaipur Jul 2018 Aug 2021

Certifications:

- Microsoft Certified: Azure AI Fundamentals
- Introduction to Responsible AI
- Complete Generative AI Course with LangChain and HuggingFace